

Annual IEEE Symposium on Foundations of Computer Science 2024

Matej Belšak^{2*}, Nik Čadež², Klemen Kavčič², Tomaž J. Leonardis¹,
Bor Pangeršič¹, Marko Rozman², Pia Sotlar², Vanja Stojanović¹

¹Faculty of Mathematics and Physics, University of Ljubljana

²Faculty of Computer and Information Science, University of Ljubljana

Abstract

We describe and explore the IEEE Annual Symposium on Foundations of Computer Science (FOCS) conference, which covers a wide area of theoretical computer science and mathematical foundations. We shortly describe 6 proceedings from the conference exploring new breakthroughs in areas such as graph theory, cryptography, compression algorithms etc.

Introduction

The FOCS conference or IEEE Annual Symposium on Foundations of Computer Science is an academic conference that covers a broad range of theoretical computer science. It is sponsored by the IEEE Computer Science Technical Committee on the Mathematical Foundations of Computing (TCMF) [1].

The FOCS 2024 took place in Chicago - Voco Chicago Downtown, from October 27-30, 2024 [2]. It covered a variety of topics, for submissions the following were mentioned [2]:

- Communication complexity
- Circuit complexity
- Average-case algorithms and complexity
- High-dimensional algorithms
- Online algorithms
- Parameterized algorithms
- Spectral methods
- Streaming algorithms
- Randomized algorithms
- Cryptography
- Computational complexity
- Algorithms and data structures
- Quantum computing
- Foundations of machine learning
- Algorithmic coding theory
- Sublinear algorithms
- Algorithmic graph theory
- Continuous optimization
- Foundations of fairness, privacy and databases
- Pseudorandomness and derandomization
- Markov chains
- Analysis of Boolean functions
- Economics and computation
- Combinatorial optimization

- Algebraic computation
- Approximation algorithms
- Parallel and distributed algorithms
- Computational learning theory
- Computational geometry
- Algorithmic game theory
- Combinatorics

The official welcome message of FOCS 2024 states that nearly 500 papers were submitted, but does not specify the exact number, out of which 133 were accepted and 131 were presented as talks during the event.

In the following sections we present curated papers from the conference.

O(1) Insertion for Random Walk d-ary Cuckoo Hashing up to the Load Threshold

Random walk d-ary cuckoo hashing

Cuckoo hashing algorithm's basic idea is to resolve collisions by using two hash functions instead of one. When inserting a new object x into the table, if the slot $h_1(x)$ is occupied, the existing object x' is replaced by x , and then x' is inserted into slot $h_2(x')$. If the number of iterations exceeds a threshold, the whole table is rehashed with new hash functions. Random Walk d -ary Cuckoo Hashing generalizes the idea by using d hash functions and using a random walk to choose the next hash function in case of a collision. Standard cuckoo hashing, equivalent to the $d = 2$ case, has a load threshold of 0.5, meaning it can use up to 50% of the hash table space. The d -ary hashing improves this threshold. For example, the $d = 3$ case has the threshold at approximately 0.9, while the insertion time increases linearly with

*Corresponding author: mb4390@student.uni-lj.si
Ljubljana, March 2025

d . The insertion algorithm guarantees $\mathcal{O}(1)$ lookup and deletion time, as the object can be retrieved by checking its d positions.

Insertion time

This paper shows that for any $d \geq 4$ hashes and load factor $c < c * d$, the expectation of the random walk insertion time is constant. It shows that the expected number of reassignments during insertion does not depend on the size of the hash table m , but only on d and c . The article uses bipartite graphs as a representation for the hash functions and objects. In the graph, the objects in set X are connected to their d possible locations in the hash table Y . In this representation, a valid perfect matching corresponds to a valid assignment of objects to slots. The existence of such a matching is subject to Hall's Theorem, which states that a perfect matching exists if and only if every subset $W \subseteq X$ is smaller than its neighborhood in Y . Neighborhood meaning all nodes connected to at least one vertex in W . The paper shows that when the load factor is below the threshold $c * d$, such perfect matchings exist with high probability as the graph exhibits strong expansion properties. The article identifies "bad" sets which have few connections to the rest of the graph. In these sets, the walk might get stuck, but the authors prove that the random walk is unlikely to hit such a set in the first $\mathcal{O}(i^{999})$ steps, and that any random walk which avoids it for the first $\mathcal{O}(i^{999})$ steps is likely to finish in $\mathcal{O}(i)$ steps.

Super-polynomial tail bounds

The paper also provides super-polynomial tail bounds on the insertion time, showing that the probability of a walk exceeding ℓ steps decays exponentially in ℓ , with the exponent approaching 1 as d increases. They relate their findings to previous work on BFS-based insertion algorithms, noting that random walk insertion achieves comparable or better performance without the computational overhead of BFS path searches [3].

Computing the 3-Edge-Connected Components of Directed Graphs in Linear Time

Abstract

The paper describes a significant improvement of a randomized (Monte-Carlo) algorithm for computing the 3-edge-connected components of a digraph with m edges in polylogarithmic time $\tilde{O}(m^{3/2})$. The al-

gorithm described bests the previous one by being deterministic and computable in linear time.

Preliminaries and primary problem

This algorithm solves the problem of finding **3-edge-connected components in directed graphs**, for preliminaries; Let $G = (V, E)$ be a strongly connected directed graph with $|V(G)| = n$ and $|E(G)| = m$. Generally a set of edges $C \subseteq E$ is a **cut** if $G \setminus C$ is not strongly connected i.e. there does not exist a directed path between every pair of vertices, if $|C| = k$ we refer to C as a **k -sized cut** of G . Hence a digraph G is **k -edge-connected** if it has no $(k - 1)$ cuts.

We say that two vertices v and w are **k -edge-connected**, and we denote this relation by $v \leftrightarrow_k w$, if there are k -edge-disjoint directed paths from v to w and k -edge-disjoint directed paths from w to v . We define a **k -edge-connected component** of a digraph G as a maximal subset $U \subseteq V(G)$ such that $u \leftrightarrow_k v, \forall u, v \in U$.

How they achieved this improvement

The new method relies on a substructure of digraphs, known as **2-connectivity-light graph** (denoted 2CLG). This is because the decomposition of digraphs into a collection of 2CLGs exists in linear time, and maintains the 3-edge-connected components of the original graph. They also define the minimal 2-in and -out sets, which contain vertices with out- or in-degree of 2. Formally, we define both here.

Definition 1. A **2-connectivity-light** graph G is a strongly connected digraph that contains two types of vertices; **ordinary** and **auxiliary**, that satisfy the following conditions:

1. Any two ordinary vertices are 2-edge-connected,
2. each auxiliary vertex has an in- or out-degree of 1,
3. for every vertex u with out-degree > 1 and every vertex v with in-degree > 1 , there are 2 edge-disjoint paths from u to v ,
4. for each auxiliary vertex v with out-degree (resp. in-degree) of one, all paths from v to any vertex in G (resp. from any vertex in G to v), we have exactly one common edge, the unique out-edge (resp. in-edge).

Definition 2. For a strongly connected digraph G we arbitrarily choose a start vertex s . For any vertex $v \neq s$ we define $M(v)$ as a **minimal 2-in set** that contains v , i.e. a minimal set of vertices which contains v , does not contains s , and has two incoming edges from $V(G) \setminus M(v)$, we denote by $M_R(v)$ the analogous sets in G^R , which is the reverse graph of G .

The technique is then based on the following proposition and theorem.

Proposition I.3. *Let G be a 2GLC with a fixed ordinary start vertex s . Then for any two ordinary vertices u and v , we have $v \leftrightarrow_k u$ if and only if $M(u) = M(v)$ and $M_R(u) = M_R(v)$.*

Theorem II.5. *Let G be a strongly connected digraph. In linear time, we can construct a collection H_1, \dots, H_t of 2CLG graphs, such that:*

- *For every vertex of G there is exactly one graph among H_1, \dots, H_t , that contains it as an ordinary vertex.*
- *Every two vertices u and v of G are 3-edge-connected if and only if there is an $i \in \{1, \dots, t\}$ such that u and v are 3-edge-connected.*

Thus concluding the paper [4].

Verifying Groups in Linear Time

Given an $n \times n$ multiplication table, the task is to decide whether it is a Cayley multiplication table of a group. To determine whether a multiplication table describes a group, you need to test for the identity element, inverses, and associativity. The first two can be trivially checked in $\mathcal{O}(n^2)$ time. The main challenge lies in efficiently testing associativity. Light's associativity test shows that it suffices to verify associativity for all triples (a, b, c) where $a, c \in G$ and $b \in R$, with R being a generating set for G . A deterministic algorithm can compute a generating set R of size $\lfloor \log n \rfloor$ in $\mathcal{O}(n^2)$ time, leading to an overall $\mathcal{O}(n^2 \log n)$ algorithm for associativity testing.

The paper improves this to an optimal $\mathcal{O}(n^2)$ algorithm by introducing a novel approach based on 4-associativity over a basis. The key idea is to reduce the problem to verifying a stronger associativity condition on a carefully chosen subset $S \subseteq G$ of size $\mathcal{O}(\sqrt{n})$, called a *basis*. If 4-associativity holds for S , then the entire multiplication table is associative. Here a basis for (G, \cdot) is a set $S \subset G$ such that $|S| = \mathcal{O}(\sqrt{n})$ and $S \cdot S = G$ (every element of G can be written as a product of two elements from S) and 4-associativity of S is defined by requiring that every $a, b, c, d \in S$ satisfy $((ab)c)d = (ab)(cd) = (a(bc))d = a((bc)d) = a(b(cd))$.

The construction of this basis relies on finding large subgroups $H \leq G$ where $|H| \geq \sqrt{|G|}$. This is achieved by reducing the problem to finding large subgroups in simple groups. The algorithm then employs specialized constructions based on the group's type (alternating, Lie type, etc.). The complete verification flows through these natural reductions: from the full multiplication table to basis testing, then to

large subgroup identification, and finally to simple group handling when needed. Each step preserves the $\mathcal{O}(n^2)$ bound.

Lempel-Ziv (LZ77) Factorization in Sublinear Time

Lempel-Ziv (LZ77) factorization is a string processing technique and the main component of most data compression algorithms, such as ZIP, PDF, and PNG. It separates a given string greedily from left to right into phrases $T = f_1 f_2 \dots f_z$, so that each phrase is either the first occurrence of a character or the longest prefix of the remaining suffix that has already appeared earlier in the text. Each phrase is encoded either as a that same character or as a pair (l, i) , where l is the length of the phrase and i is the position of its earlier occurrence. For example, for the string $T = b \cdot b \cdot a \cdot ba \cdot aba \cdot bababa \cdot ababa$, the LZ77 representation is $(0, b), (1, 1), (0, a), (2, 2), (3, 3), (6, 7), (5, 10)$.

In the RAM model, the theoretical lower bound for LZ77 factorization is $\mathcal{O}(n / \log_\sigma n)$, which matches the maximum number of phrases. Despite this, no algorithm achieved this optimal bound until recently. Earlier algorithms were based on suffix trees or suffix arrays, achieving $\mathcal{O}(n \log \sigma)$ time and $\mathcal{O}(n)$ space. Later improvements reduced space usage to the optimal $\mathcal{O}(n / \log_\sigma n)$, but the algorithm still required $\Omega(n)$ time in the worst case.

Kempa and Kociumaka introduced the first sublinear-time algorithm for LZ77 factorization, running in $\mathcal{O}(n / \sqrt{\log n})$ time for binary alphabets and $\mathcal{O}((n \log \sigma) / \sqrt{\log n})$ time for larger alphabets, while using optimal $\mathcal{O}(n / \log_\sigma n)$ space. They developed a novel index structure that can be built in sublinear time and efficiently finds the leftmost previous occurrence of a substring. This index makes the computation of the Longest Previous Factor (LPF) for each position in the string fast and answers substring occurrence queries in $\mathcal{O}(\log^\epsilon n)$ time.

Instead of relying on the classical method based on Range Minimum Queries (RMQ) over the suffix array—which is too slow for sublinear-time construction—they use a sampling-based approach. The idea is that instead of looking through all the text positions they select a small, representative sample (S) of positions in non-periodic regions (size $\mathcal{O}(n / \log_\sigma n)$). The rule is that if two substrings are equal their starting positions are either both included or both excluded from S .

To check if a substring has appeared before, the problem is turned into a special kind of Range Minimum Query (RMQ). Prefix RMQ finds the smallest

value in a range, but only if the prefix matches a given pattern. The authors create a fast and memory-efficient data structure to perform these prefix RMQ queries.

For periodic regions, the algorithm uses two different approaches. If the pattern is only partially periodic, it uses sorted runs and special range queries (3-sided RMQ) to handle them. If the pattern is fully periodic, the algorithm marks the starting positions with a bitvector (a sequence of 0 and 1, where 1 defines the starting positions of the pattern). It then uses fast rank/select queries on this bitvector to quickly find the relevant patterns.

So by using this sampling-based approach with prefix RMQ structure, the authors achieve an index that maintains optimal $O(n/\log_\sigma n)$ space complexity, sublinear preprocessing time, and efficient query performance [5].

Tight bounds for Classical Open Addressing

The authors investigate the optimal tradeoff between time and space in classical open addressed hash tables with a high load factor of $1 - \epsilon$. They introduce the *Rainbow Hash Table*, which achieves $O(1)$ expected query time and $O(\log \log(1/\epsilon))$ update time, while maintaining near-full hash table. They prove this tradeoff is optimal, demonstrating that no open-addressed hash table can achieve better time complexity while supporting such a high load factor.

Paper first introduced *Rainbow Cell*, consisting of $n^{1/4}$ buckets, each containing $n^{3/4}$ slots of which $n^{1/2}$ slots are special *sky slots*. *Rainbow Cell* is specialized hash table that operates at load factor of 1 with $O(1)$ time complexity for insertions and deletions, while queries have $O(n^{3/4})$ time complexity. Each element is randomly assigned either *heavy* (high probability) or *light* (low probability) state with status hash function. Heavy keys are stored only in bucket that it was hashed to. Light keys are stored only in sky slots and can be moved between buckets without disruption structure.

The idea is to reduce problem to subproblems and *Rainbow Hash Table* is introduced. It is a tree structure where each subproblem is implemented as a *Rainbow Cell*. Elements are randomly assigned a color, which is used to help locate level of where element is going to end up. The probability distribution of colors ensures that most elements are placed in lower levels. Meaning, at lower levels of tree, subproblems become significantly smaller and tree becomes wider.

Next, *Rainbow Hash Table* is extended to support dynamic resizing, meaning the hash table could in-

crease and decrease in size while always remaining full (load factor of 1) and preserving its efficiency. Expected query time is $O(1)$ and expected update time of $O(\log \log n)$, where n is current size of hash table.

The final algorithm supports a load factor up to $1 - \epsilon$ and ensures insertions and deletions remain efficient. The last sections focus on their main result, proving the optimal time complexity for insertions and deletions in classical open addressing with a high load factor of $1 - \epsilon$ is $O(\log \log(1/\epsilon))$, while queries can be performed in $O(1)$. This result resolves the open question of whether it is possible to achieve operations with $O(1/\epsilon)$ time complexity for insertions, deletions, and queries simultaneously in an open-addressed hash table [6].

Minor Containment and Disjoint Paths in almost-linear time

Abstract

The paper *Minor Containment and Disjoint Paths in Almost-Linear Time* by Tuukka Korhonen, Michał Pilipczuk, and Giannos Stamoulis introduces an advanced algorithm that determines whether a graph H is a minor of another graph G in $\mathcal{O}_H(n^{1+o(1)})$ time, where n represents the number of vertices in G . This algorithm is a further development on the work previously done by Kawarabayashi, Kobayashi, and Reed which had complexity quadratic relative to the number n of vertices of graph G .

Minor containment problem

In graph theory, a graph H is considered a minor of a graph G if H can be derived from G through a series of vertex deletions, edge deletions, and edge contractions. The ability to test for minor containment efficiently is crucial, as it underpins the recognition of minor-closed graph classes—collections of graphs where any minor of a member graph is also included in the class. The Graph Minor Theorem (Robertson–Seymour) asserts that such classes can be characterized by a finite set of forbidden minors, making efficient minor testing algorithms highly valuable.

Beyond unrooted minor containment, the authors also address the rooted version of the problem. Here, G is provided with a set of roots $X \subseteq V(G)$, and certain branch sets of the desired minor model of H must include specified subsets of X . The proposed algorithm solves this in $\mathcal{O}_{H,|X|}(m^{1+o(1)})$ time, where m denotes the total number of vertices and edges

in G . This encompasses the Disjoint Paths problem, a classic challenge in graph algorithms that seeks paths between specified pairs of terminals without shared vertices or edges. The new algorithm achieves this in $\mathcal{O}_k(m^{1+o(1)})$ time, with k being the number of terminal pairs.

Main Contributions

The efficiency of the proposed algorithm of the authors comes from two primary components:

- **Implementation of the Irrelevant Vertex Technique:** Utilizing the *dynamic treewidth* data structure developed by Korhonen et al., the algorithm applies the irrelevant vertex technique of Robertson and Seymour in almost-linear time on apex-minor-free graphs. This technique identifies and removes vertices that do not contribute to the minor structure, simplifying the graph without affecting the minor relationship.
- **Application of Recursive Understanding:** Leveraging recent advancements in almost-linear time flow/cut algorithms, the recursive understanding technique effectively reduces complex graphs to apex-minor-free graphs. This reduction simplifies the problem, making it more tractable for the minor containment test.

Conclusion

These innovations collectively enable the algorithm to perform minor containment and solve the disjoint paths problem with better efficiency. In the paper the authors provide detailed proofs for theoretical advancements and complexity analyses.

The authors also believe that the key new insights of the paper, namely - the almost-linear time implementation of recursive understanding and the almost-linear time implementation of the irrelevant vertex rule on apex-minor-free graphs with the dynamic treewidth data structure could find applicability much more broadly in the context of fixed-parameter algorithms and computational problems in the theory of graph minors [7].

References

- [1] IEEE Symposium on Foundations of Computer Science, *Ieee focs conference*, Accessed: 2025-02-27, 2024. [Online]. Available: <https://ieeefocs.org>.
- [2] IEEE Symposium on Foundations of Computer Science, *Focs 2024 - 65th annual ieee symposium on foundations of computer science*, Accessed: 2025-02-27, 2024. [Online]. Available: <https://focs.computer.org/2024/>.
- [3] T. Bell and A. Frieze, "O(1) insertion for random walk d-ary cuckoo hashing up to the load threshold," in *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, 2024, pp. 106–119. doi: 10.1109/FOCS61266.2024.00017.
- [4] L. Georgiadis, G. F. Italiano, and E. Kosinas, "Computing the 3-edge-connected components of directed graphs in linear time," in *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, 2024, pp. 62–85. doi: 10.1109/FOCS61266.2024.00015.
- [5] D. Kempa and T. Kociumaka, "Lempel-ziv (lz77) factorization in sublinear time," in *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, 2024, pp. 2045–2055. doi: 10.1109/FOCS61266.2024.00122.
- [6] M. A. Bender, W. Kuszmaul, and R. Zhou, "Tight bounds for classical open addressing," in *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, 2024, pp. 636–657. doi: 10.1109/FOCS61266.2024.00047.
- [7] T. Korhonen, M. Pilipczuk, and G. Stamoulis, "Minor containment and disjoint paths in almost-linear time," in *2024 IEEE 65th Annual Symposium on Foundations of Computer Science (FOCS)*, 2024, pp. 53–61. doi: 10.1109/FOCS61266.2024.00014.